

# Hive HQL Hands-on: Logs & Errors Analysis

## Create External Table for Log Data

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hive  
Logging initialized using configuration in jar:file:/usr/lib/hive/lib/hive-comm  
n-0.10.0-cdh4.4.0.jar!/hive-log4j.properties  
Hive history file=/tmp/cloudera/hive_job_log_b8404494-8a73-4bc8-abf0-9b186be62e  
0_248843411.txt  
hive> CREATE EXTERNAL TABLE IF NOT EXISTS logs (  
log_date STRING,  
log_level STRING,  
service_name STRING,  
message STRING  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION '/user/hive/logs/';  
OK  
Time taken: 2.073 seconds  
hive>
```

## View All Logs

```
hive> select * from logs;  
OK  
2025-05-01 11:39:00 INFO InventoryService Order placed  
2025-05-01 11:25:00 WARN PaymentService Address not verified  
2025-05-01 12:49:00 ERROR ShippingService Shipping error  
2025-05-01 14:33:00 WARN ShippingService Slow response time  
2025-05-01 10:27:00 INFO ShippingService Inventory updated  
2025-05-01 14:11:00 INFO PaymentService Order placed  
2025-05-01 12:28:00 ERROR OrderService Payment failed  
2025-05-01 13:54:00 ERROR OrderService Payment failed  
2025-05-01 14:54:00 ERROR AuthService Order placement failed  
2025-05-01 10:54:00 WARN OrderService Address not verified  
2025-05-01 14:38:00 ERROR ShippingService Payment failed  
2025-05-01 11:18:00 INFO OrderService Order placed  
2025-05-01 12:44:00 INFO PaymentService Order placed
```

## Count Total Logs

```
2025-05-17 06:43:56,167 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.66 :  
c  
2025-05-17 06:43:57,185 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7  
sec  
2025-05-17 06:43:58,199 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7  
sec  
2025-05-17 06:43:59,214 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7  
sec  
2025-05-17 06:44:00,222 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7  
sec  
2025-05-17 06:44:01,230 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7  
sec  
MapReduce Total cumulative CPU time: 2 seconds 770 msec  
Ended Job = job_202505170500_0001  
MapReduce Jobs Launched:  
Job 0: Map: 1 Reduce: 1 Cumulative CPU: 2.77 sec HDFS Read: 6357 HDFS Wri  
: 4 SUCCESS  
Total MapReduce CPU Time Spent: 2 seconds 770 msec  
OK  
100  
Time taken: 26.837 seconds  
hive>
```

## Count Logs by Log Level

```
2025-05-01 11:39:00,207 Stage: 2 map = 1000, reduce = 1000, cumulative CPU 2.74 sec
MapReduce Total cumulative CPU time: 2 seconds 740 msec
Ended Job = job_202505170500_0002
MapReduce Jobs Launched:
Job 0: Map: 1 Reduce: 1 Cumulative CPU: 2.74 sec HDFS Read: 6357 HDFS Write: 25 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 740 msec
OK
ERROR      38
INFO       33
WARN       29
Time taken: 20.891 seconds
hive>
```

## Count ERROR logs per service

```
Ended Job = job_202505170500_0002
MapReduce Jobs Launched:
Job 0: Map: 1 Reduce: 1 Cumulative CPU: 1.85 sec HDFS Read: 6357 HDFS Write: 85 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 850 msec
OK
AuthService      4
InventoryService  5
OrderService     10
PaymentService   8
ShippingService  11
Time taken: 20.008 seconds
hive>
```

## Filter Logs by Date

```
Total MapReduce CPU Time Spent: 920 msec
OK
2025-05-01 11:39:00 INFO InventoryService Order placed
2025-05-01 11:25:00 WARN PaymentService Address not verified
2025-05-01 12:49:00 ERROR ShippingService Shipping error
2025-05-01 14:33:00 WARN ShippingService Slow response time
2025-05-01 10:27:00 INFO ShippingService Inventory updated
2025-05-01 14:11:00 INFO PaymentService Order placed
2025-05-01 12:28:00 ERROR OrderService Payment failed
2025-05-01 13:54:00 ERROR OrderService Payment failed
2025-05-01 14:54:00 ERROR AuthService Order placement failed
2025-05-01 10:54:00 WARN OrderService Address not verified
2025-05-01 14:38:00 ERROR ShippingService Payment failed
2025-05-01 11:18:00 INFO OrderService Order placed
2025-05-01 12:44:00 INFO PaymentService Order placed
```

## Create Partitioned Table by Date

```
hive> CREATE EXTERNAL TABLE logs_partitioned (  
  log_time STRING,  
  log_level STRING,  
  service_name STRING,  
  message STRING  
)  
PARTITIONED BY (log_date STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;  
OK  
Time taken: 0.839 seconds
```

## Create a View for ERROR Logs

```
total mapreduce job time spent: 0.00 msec  
OK  
2025-05-01 12:49:00 ERROR ShippingService Shipping error  
2025-05-01 12:28:00 ERROR OrderService Payment failed  
2025-05-01 13:54:00 ERROR OrderService Payment failed  
2025-05-01 14:54:00 ERROR AuthService Order placement failed  
2025-05-01 14:38:00 ERROR ShippingService Payment failed  
2025-05-01 10:51:00 ERROR ShippingService Order placement failed  
2025-05-01 11:28:00 ERROR PaymentService Order placement failed  
2025-05-01 12:04:00 ERROR ShippingService User authentication failed for u  
ser 105  
2025-05-01 13:56:00 ERROR OrderService Order placement failed  
2025-05-01 13:32:00 ERROR InventoryService Shipping error  
2025-05-01 11:27:00 ERROR InventoryService Inventory sync error  
2025-05-01 13:03:00 ERROR ShippingService Shipping error  
2025-05-01 10:28:00 ERROR OrderService Order placement failed  
2025-05-01 10:51:00 ERROR InventoryService Inventory sync error  
2025-05-01 12:41:00 ERROR OrderService Shipping error  
2025-05-01 14:49:00 ERROR ShippingService User authentication failed for u  
ser 106
```